

Facilitator guide — Prompting Techniques

Facilitator guide · Prompting Techniques · 6 activities · ~115 min

Purpose

This module develops a single practical skill: instructing the AI Assistant clearly enough to get a useful result. Participants progress from writing one well-formed prompt, through exploring how the AI responds to small changes, to refining and verifying outputs. The skills established here are assumed by every later AI-assisted activity in the workshop — building visualisations, slide decks, and disruption reports — so the module functions as the foundation for that work.

By the end, a participant should be able to: write a prompt that states its objective, audience, scope, and format; adjust a prompt deliberately and observe the effect; choose between an iterative conversation and a single structured prompt; and apply a consistent check to an AI-generated draft before using it.

Session at a glance

#	Activity	Time	Format
1	Build a clear prompt	~15 min	Individual, then share
2	Explore with the AI Assistant	~20 min	Pairs or small teams
3	Iterative vs single prompt	~30 min	Pairs or small teams
4	Refine your prompt	~20 min	Individual (self-paced)
5	Use a previous report as a template	~10 min	Pairs or small teams
6	Verify the AI's output	~20 min	Individual or pairs

Running the session

Preparation. Confirm before the session that each table can sign in to the platform and open the AI Assistant within its country project. Generation takes up to a minute per request, so a reliable connection matters more here than in most modules. Prepare one prompt of your own in advance so you can demonstrate from a known-good example rather than improvising.

How to demonstrate. Most activities are hands-on. For each, show the first step on the shared screen — enough that participants recognise the interface — then let them work from the handout at their own pace. Resist working through a whole activity on screen: participants learn prompting by writing prompts, not by watching you write them. The on-screen moments that matter most are flagged under **Demonstrate** in each activity below.

Grouping. Activities 2, 3 and 5 work well in pairs or small teams, where participants compare outputs and discuss why they differ. Activities 1 and 4 are individual; participants need to form their own judgement before comparing notes.

Pacing. Activity 1 is foundational and should not be abbreviated, even when time is short — the later activities assume participants can name the six dimensions of a good prompt. If the session runs behind, shorten Activity 5 (the shortest and most self-explanatory) rather than Activity 1 or 6.

The message to carry through. A strong prompt is rarely the first one written; refining is the normal way of working, not a sign of failure. Reinforce this at each activity so participants do not become discouraged when a first attempt returns a generic answer.

The activities

1. Build a clear prompt · ~15 min · individual

What happens. Participants take a task they do regularly and, before writing anything, check it against six dimensions: objective, audience, geography/period/scope, interpretation guidance, output format, and guardrails. They then write a single prompt that addresses each, and share it with one neighbour.

Demonstrate. Put a deliberately vague prompt on screen — *"tell me about our data"* — and ask the group what the AI is likely to return. Then show a version that names objective, audience, scope and format, and contrast the two. The point lands faster from a side-by-side than from explanation.

Say something like. *"The six dimensions are not a form to fill in. They are a 30-second check before you type, so your first message is already close."*

What a good result looks like. A prompt — a short paragraph is enough — in which a reader can predict what the AI will return, because the objective, audience, scope and format are all stated.

Watch for.

- Aiming for one perfect prompt. Reassure them the six dimensions just produce a stronger *first* message; refining still follows.
- A vague prompt. Rather than rewriting it, ask which of the six dimensions is missing and have them add it.
- Guardrails skipped — the dimension most often dropped. Remind them to tell the AI to stay within the data shown and flag uncertainty.

2. Explore with the AI Assistant · ~20 min · pairs

What happens. Participants run a supplied base prompt (*ANC1 utilization by region over the last 12 months*), re-run it changing exactly one element — indicator, period, or geographic level — and observe how the output changes. They ask one follow-up to see the exchange can continue as a conversation, and save anything worth keeping to their personal folder.

Demonstrate. Paste the base prompt into the AI Assistant on screen and run it. Then change a single word — swap the indicator — and run it again, so the group sees one change produce one difference before they try it themselves.

Say something like. *"Change one thing, look at what moved, then change the next. If you change two at once you won't know which one did it."*

What a good result looks like. Two or three outputs that differ in a single, identifiable way, and a participant who can say which change produced which difference.

Watch for.

- Changing two things at once — the most common error. Hold them to one variable per run.
- Useful output left unsaved. Prompt them to save a good chart or interpretation to their folder before moving on.

3. Iterative vs single prompt · ~30 min · pairs

What happens. Participants build the same three-slide data quality report twice. In Exercise A they reach it through four prompts sent in sequence, reacting to each response. In Exercise B they open a fresh conversation and reach it with one structured prompt. They compare the two outputs and complete a short reflection, supported by a "when to use each" table.

Demonstrate. Before they start Exercise B, show how to open a *fresh* conversation, and explain why it matters: if they continue in the Exercise A thread, the single prompt inherits all that earlier context and the comparison is meaningless.

Say something like. *"Neither approach is better. Iterate when you're still working out what you want; use one structured prompt when you already know and will repeat it."*

What a good result looks like. Two comparable reports and a participant who can state the trade-off: the iterative path gives more control and suits exploration; the single prompt is faster and suits repeatable, routine outputs.

Watch for.

- Declaring one approach "better". Both use the same data and methods — the choice depends on the task.
- Exercise B run in the same thread as A. Check they started a new conversation, or the result is contaminated.

4. Refine your prompt · ~20 min · individual (self-paced)

What happens. Participants choose one topic — outliers, completeness, or disruptions — and run three rounds on it, each in a new conversation: a simple request, the same request with context added, then a version asking for prioritisation and next steps. After each round they note what improved.

Demonstrate. Walk through the three-round table once with a neutral topic so the progression is clear — round 1 bare, round 2 adds *where/when*, round 3 adds *why/who* — then let them run their own topic.

Say something like. *"Don't just notice that the last one was better — name what you added that made it better. That's the part you can reuse."*

What a good result looks like. Three visibly different slides on one topic, and a participant who can name what each round added — round two typically adds *where* and *when*, round three adds *why* and *for whom*.

Watch for.

- Running rounds as follow-ups in one thread instead of fresh conversations — they then can't see each prompt's standalone effect.
- Seeing improvement without seeing the cause. Push them to articulate the specific addition each round, and to save the strongest phrasing.

5. Use a previous report as a template · ~10 min · pairs

What happens. Rather than describing a structure in words, participants show the AI an example. They upload a report they already trust to Assets, include it in a fresh conversation through the three-dot menu, and ask the AI to reproduce its structure for a new period or scope.

Demonstrate. This activity turns on one fiddly sequence — show it slowly on screen: **Assets** → **Upload assets** → **select PDF**, then in a conversation **three-dot menu** → **Include file** → **select it**. Point out that the menu must show the file as *attached*.

Say something like. *"Uploading a file is not the same as including it. If the AI says it can't see your report, this is almost always why."*

What a good result looks like. A new report that follows the structure of the uploaded model, checked section by section against it.

Watch for.

- File uploaded but not included in the conversation — the single most common failure. Confirm the three-dot menu shows it attached.
- The AI drifting from the template — simplifying a chart, dropping a section. Have them compare output to model side by side before using it.
- A PDF too large to upload. Suggest splitting it into chapters or extracting the relevant pages.

6. Verify the AI's output · ~20 min · individual or pairs

What happens. Participants apply a repeatable check to an AI-generated draft: read once and mark every factual claim, sort claims by risk, verify high-risk claims manually against the data, ask the AI to quote its source for medium-risk claims, and finish with a short consistency scan.

Demonstrate. Take a short AI draft on screen and mark two or three claims aloud — a number, a cause-and-effect statement — then show the risk table and place those claims into it. Modelling the marking step removes most of the hesitation.

Say something like. *"The AI writes fluently, which makes wrong numbers sound right. Fluent is not the same as correct — every figure has to trace back to the data."*

What a good result looks like. A draft in which every retained number can be traced to its source, and any claim the AI could not support has been removed or rewritten.

Watch for.

- Rushing — the most common failure. It needs the full twenty minutes; protect the time.
- Trusting confident phrasing. Flag the tell-tale signs: round numbers, precise figures with no source, plausible numbers that can't be traced.
- The principle to leave them with: the person who submits the work owns every figure in it — a draft is not finished until they would put their name to it.

Closing the session

Close on verification. Prompting well produces a fast first draft; verifying well is what makes that draft the participant's own, defensible work. This pairing — generate, then verify — is the habit the rest of the workshop relies on.