

Verifying an AI-generated report

Reference · Standard FASTR Reports · ~20 min

WHY THIS MATTERS

The AI Assistant can generate a complete disruption report in 5–10 minutes. **You** sign off on it. The AI can sound confident and still be wrong — about a number, a region name, an indicator grouping, a comparison.

THE CORE RULE

Every claim in the report must be checkable against the underlying chart or table. If you can't trace a sentence back to data, flag it. Verification is what you do **before** the report leaves the team.

How a FASTR report is built — the prompts

The report builder uses a set of prompts from the prompt library:

- **Prompt 1 — FASTR disruption report.** The base report. The AI asks for country / period / subtitle, finds the available indicators, proposes groupings, you confirm, and it builds the report slide by slide.
- **Prompt 2 — Regional disruption analysis.** *Optional.* Adds one slide per subnational area.
- **Prompt 3 — Data quality assessment.** *Optional.* Adds an appendix on completeness, outliers, consistency.
- **Prompt 5 — Review slide deck.** The **verification prompt**. Run it after the report is generated — it checks the whole deck for you.

Two passes to verify

Verification has two passes, and you need **both**:

1. **Pass 1 — AI self-review.** Run **Prompt 5** so the AI checks its own output against the data.
2. **Pass 2 — Team review.** The AI can't catch everything. The team runs the manual checklist.

Pass 1 — Run the AI review (Prompt 5)

Once Prompt 1 (and any of Prompt 2 / 3) has finished generating, open the prompt library and run **Prompt 5: Review slide deck**. The AI goes through **every slide in one pass** and checks six things:

#	What Prompt 5 checks
1	Data accuracy — every number in the text matches the underlying data; unverifiable stats get flagged <code>[UNVERIFIED]</code> ; watches for hedged fabrication ("approximately X" hiding an invented figure)
2	Indicator names & direction — names match the exact platform labels; service indicators up = good, mortality indicators up = bad, not mixed up
3	Acronyms & methodology — any acronym expansion or methodology description is verified against the official docs
4	Language & framing — no causal claims, appropriate hedging, no overgeneralisation, correct health terms
5	Consistency across slides — same indicator shows the same value everywhere; names and periods consistent
6	Word count — each text block within the target range

What you get back: a summary — "*[X] slides reviewed, [Y] issues found across [Z] slides*" — with a suggested fix per issue. You choose: fix all automatically, review one by one, or done.

Pass 1 is fast and catches mechanical errors. But it only checks the deck against itself and the data. It is not the end of verification — it's the start.

Pass 2 — Team review

Prompt 5 cannot know things only your country team knows. **It cannot tell:**

- whether the indicator **groupings make sense** for your country
- whether a "disruption" is **real or explained** by a known local event (a strike, a campaign, a stockout)
- whether the **so-what is realistic** and actionable in your context
- whether **local context** is missing

That's what the manual checklist below is for. Work through it **as a team**, with the source charts open on the platform next to the report.

Before you start the team review:

- ☐ Prompt 5 has been run and its flagged issues fixed
- ☐ Your country team is together — this is a team activity, not solo
- ☐ The source charts are open on the platform, next to the report

Verify **Section A + Section B** for every report. Add **Section C / D** only if you ran Prompt 2 / Prompt 3.

Section A — General checks (every report)

☐	Check	Notes
☐	All slides generated — none missing, none half-rendered	
☐	The country name is correct and consistent on every slide	
☐	The analysis period is correct on every chart	
☐	No placeholder text left behind — e.g. [COUNTRY] , [VERIFY] , [UNVERIFIED]	
☐	The last page (FASTR link / closing slide) is genuinely the last slide	

Section B — Prompt 1: base disruption report

☐	Check	Notes
☐	Cover page — country, subtitle, and date all correct	
☐	The indicator groupings match what you confirmed in the chat	
☐	Each group has its own analysis slide — none merged, none dropped	
☐	Charts load and show data — no blank or broken charts	
☐	Each interpretation matches the chart on that slide	
☐	Titles are analytical, not just indicator names — "ANC1 fell 12% in the North", not "ANC1 results"	
☐	Does the so-what make sense for your country? Is it actionable?	

If a check fails: for placeholders or structure, ask the AI to regenerate the named slide. For a wrong interpretation: *"Slide 6 says volumes rose — the chart shows a fall. Please re-check."*

Section C — Prompt 2: regional disruption analysis

Only if you ran Prompt 2.

☐	Check	Notes
☐	The summary includes all subnational areas (count them)	
☐	One slide per area — none missing, none duplicated	
☐	Area names spelled correctly (matches official naming)	
☐	Each interpretation references the right area	
☐	Observed-vs-expected charts use the same scale across slides	
☐	Magnitudes match the source chart (no inflated drops, no missed spikes)	

Section D — Prompt 3: data quality appendix

Only if you ran Prompt 3.

☐	Check	Notes
☐	The appendix is numbered correctly (Annex 1 or 2)	
☐	All three DQ dimensions present: completeness, outliers, consistency	
☐	The color code is logical : green = good, red = alert	
☐	Low-quality areas or indicators are explicitly flagged in the text	
☐	DQ scores match the platform — open the DQA module side-by-side	

Final team read-through

Once every section passes, **read every interpretation together as a team** — each team member taking a section, the rest following along. Watch for:

- **Anything that reads wrong** — a number too round, a trend that's not in the chart
- **Anything exaggerated** — "dramatic", "alarming", "unprecedented" without evidence
- **Anything generic** — sentences that could apply to any country, not yours
- **Missing local context** — the strike, the supply-chain issue, the new policy. You add it.

If anyone on the team hesitates on a sentence, **flag it**. Hesitation usually means something is off.

Sign-off

Before the report leaves the team, the team must be able to say:

"Prompt 5 has been run and its issues fixed. Every number is checkable. Every interpretation reflects what we know. No sentence is generic. No claim is unsupported."

If you can't say that yet, keep iterating.

What's next

Once the report passes both verification passes, finalize the formatting (titles, page numbers, logos), export, and disseminate using your country action plan.